

Analysis of Algorithms for the Euclidean Traveling Salesman Problem

Iago Zagnoli Albergaria¹, Manuel Junio Ferraz Cardoso¹

¹Departamento de Computação – Universidade Federal de Minas Gerais (UFMG)

{iagozag,mjfc}@ufmg.br

Abstract. *The Euclidean Traveling Salesman Problem (ETSP) seeks the shortest tour visiting a set of points in two-dimensional Euclidean space. This paper examines branch-and-bound and two approximation algorithms, focusing on their complexity and performance when solving the problem. Furthermore, the experimental comparisons highlight trade-offs between solution accuracy and computational efficiency / memory consumption.*

1. Introduction

The Euclidean Traveling Salesman Problem (ETSP) is a classical optimization problem that involves finding the shortest possible path that visits a set of points in a two-dimensional Euclidean space. It is an instance of the Traveling Salesman Problem (TSP) that has been widely studied due to its computational complexity and extensive applications in logistics, transportation, and circuit design. In the Euclidean version, the distance between two vertices is computed using the Euclidean distance in a two-dimensional plane.

This paper focuses on the analysis of algorithms designed to solve the ETSP, specifically with branch-and-bound and two approximation techniques. Branch-and-bound is an exact method that explores possible solutions while pruning suboptimal branches based on bounds. On the other hand, approximation algorithms provide faster results, though they may not always produce optimal solutions. The main goal of this study is to evaluate and compare these algorithms, analyzing their computational complexity, memory usage, and solution accuracy.

2. Implementation

2.1. Branch and Bound

The initial implementation of the algorithm used Best-First Search and, during testing, performed well in smaller instances. However, when scaling to larger cases, the algorithm consumed excessive memory, causing the program to crash, and, in some instances, even crashing the computer. As a consequence, there was a switch to Depth-First Search.

First, the initial bound for the graph is calculated, the optimal value is initialized to infinity, and a vector of booleans is created, with only the first vertex marked as visited. Each state node has five parameters: the lower bound of the remaining graph, the cost up to the node, the level of recursion, the current vertex, and the vector of visited vertices. Moreover, the vector of booleans is passed by reference to save memory.

The initial bound is calculated as the sum of the two smallest edges leaving each vertex, divided by two. This is a lower bound because every path uses two edges per

vertex, and selecting the two smallest edges provides the best possible estimation. When moving to a node, the bound is decreased by the sum of the second smallest edge of the vertex being left and the first smallest edge of the vertex being entered, divided by two. In addition, the value of the edge being traversed is increased.

Other than that, the function goes straight to the point. It tries every permutation of the vertices, with the exception that if the bound for a vertex is greater than a previous estimation, that node is pruned from the recursion.

2.2. Twice Around The Tree

It is known that the ETSP is a NP-Hard problem, meaning that no algorithm can find the exact solution in polynomial time. However, if we stop chasing the optimal answer and try to find another one that is close to it, we can be more efficient than that. One of the algorithms that implement this idea is the Twice Around the Tree (TATT) algorithm.

To understand how the algorithm works, we need to recognize a property of the problem: in the optimal path, removing any edge from it results in a tree that contains all the points. In other words, the path will become a Spanning Tree. We will refer to the spanning tree obtained this way as ST.

The TATT algorithm begins by generating a Minimum Spanning Tree (MST) in the given set of points. In this implementation, Prim's algorithm is used to construct the MST. Since this tree is the least weighted spanning tree, its weight will be no greater than that of the ST. By doubling the edges in the MST, it will contain an Eulerian Cycle that is, at most, twice as weighted as ST (and also the optimal path, since ST is not heavier than the optimal answer).

Subsequently, the algorithm finds the Eulerian Cycle contained in this MST (after double all its edges) using the Depth First Search (DFS) algorithm. Finally, any duplicate occurrences of points are removed from the path. For example, if x, y, z are points and y has already appeared in the path, we change the subsequence xyz to xz .

However, in this implementation the edges in the MST are not duplicated. Instead, the DFS simply inserts each point at the end of the path when it is visited during a Preorder traversal of the MST. This results in a path that includes every point exactly once, which is the answer of the algorithm (and is at most twice the optimal cost).

2.3. Christofides

Other algorithm that finds an answer that is close to the optimal is the Christofides algorithm. But it find an answer that is closer than the found by TATT algorithm.

The first part of the algorithm is the same that the first part of TATT: find the MST (we will call it T) using Prim's algorithm. After this, a subgraph G that contains all the vertices with odd degrees in T , and all the edges between these vertices (so, a complete subgraph). Since the sum of the degrees of every vertex in any graph is even, the number of vertices with odd degrees of every graph is even, so T has an even number of vertices with odd degree.

The next step is to calculate a minimum perfect matching of G (actually, a maximum matching is calculated, but we change the weight of the edges in order to turn the matching a minimum perfect matching in the original weight). Since G is a complete

graph with an even number of vertices, we know that this matching is always perfect. On this implementation, Blossom algorithm was used. Our implementation has used an open source library, LEMON, to find the matching ([Dezső et al. 2011]).

Let M be the graph that contains all the vertices of G and the edges used in the matching. We create a new graph H that contains all the vertices in the input and all the edges in T and in M . If some edge is in both of them, we add it duplicated (that is why we have used a data structure that increases the complexity of the algorithm). Now, H is a connected graph that every vertex has an even degree (so, H is an Eulerian Graph).

The last step is find this Eulerian cycle. The Hierholzer's Algorithm was used to make this. Then it is transformed in a Hamiltonian path the same way as described in the Twice Around The Tree algorithm (the idea, not the way it was implemented): remove any occurrence of any vertex that has already appeared before. The final path is the answer of the algorithm, with an cost at most 1.5 times the optimal one.

3. Results

Most of the tests were from the TSPLIB library [Reinelt 1991], available at <http://comopt.ifl.uni-heidelberg.de/software/TSPLIB95/>. Some additional tests were taken from the work of [Dejnek et al. 1997].

3.1. Branch and Bound

The branch-and-bound method always provides the optimal answer, but the asymptotic complexity is $O(N!)$. As a result, even for instances with a small number of vertices, the program takes over 30 minutes and produces NA. In fact, this happens with all tests from the TSPLIB library [Reinelt 1991]. Given that, three additional tests were added: two of them are from [Dejnek et al. 1997], and one consisting of a simple line graph with 15 points $(0,0), (0,1), \dots, (0,14)$.

The first test case is a Kalmanson point set with 10 points that executes in less than 1 millisecond. The second test case is a Demidenko point set with 15 points that is completed in 16 milliseconds. The third test case is a manually created line graph with 15 points designed for comparison with the Demidenko point set. This comparison highlights that the Demidenko point set is an easier case for the ETSP, as it completes in only 16 milliseconds, whereas the line graph, with the same number of vertices, fails to terminate even after 30 minutes of execution. The reason the first two test cases are computed quickly is explained in the article by [Dejnek et al. 1997].

3.2. Twice Around The Tree

This algorithm gives an approximated answer, not the optimal one. So, let's analyze the answer. The graphic below shows how close to the optimal cost the found was for every test case:

Average approximation to optimal cost using the Twice Around The Tree algorithm




Figure 1. Histogram of approximation rates

The algorithm has a guarantee that the found answer is at most the double of the optimal. But, as we can see, it usually gives a cost that goes from 1.3 to 1.5 times the optimal.

Let n be the number of vertices. The version of the algorithm used to find the MST is $O(n^2)$, and the DFS algorithm to find the Hamiltonian cycle is $O(n \times \log_2(n))$, since now the graph has $n - 1$ edges. This $\log_2(n)$ factor is present because of the data structure used in the adjacency list: the class *multiset* in C++, which is a binary tree, was used (the reason to this is the necessity to represent some edges twice, as needed by the Christofides algorithm). This gives a total of $O(n)$ of time complexity, which can be visualized in the graphic below.

Number of points x execution time using the Twice Around The Tree algorithm




Figure 2. Relation between the number of points and execution time of the TATT

To finish this analysis, we can check the space complexity: the original graph (the edges between each pair of points) are not created. Instead, we just save the coordinates of each point, what cost $O(n)$ space. The algorithm to generate the MST have another cost of $O(n)$ space. Since we have $n - 1$ edges in the MST, we have, also, $O(n)$ to keep the edges of the MST. To finish it, the DFS have a final $O(n)$ complexity on the worst case. So, this algorithm have a space complexity of $O(n)$.

3.3. Christofides

Let's follow the same order of the analysis of the TATT algorithm: starting from the quality of the answer. We can verify one advantage of this algorithm looking at the next graphic:

Average approximation to optimal cost using the Christofides algorithm

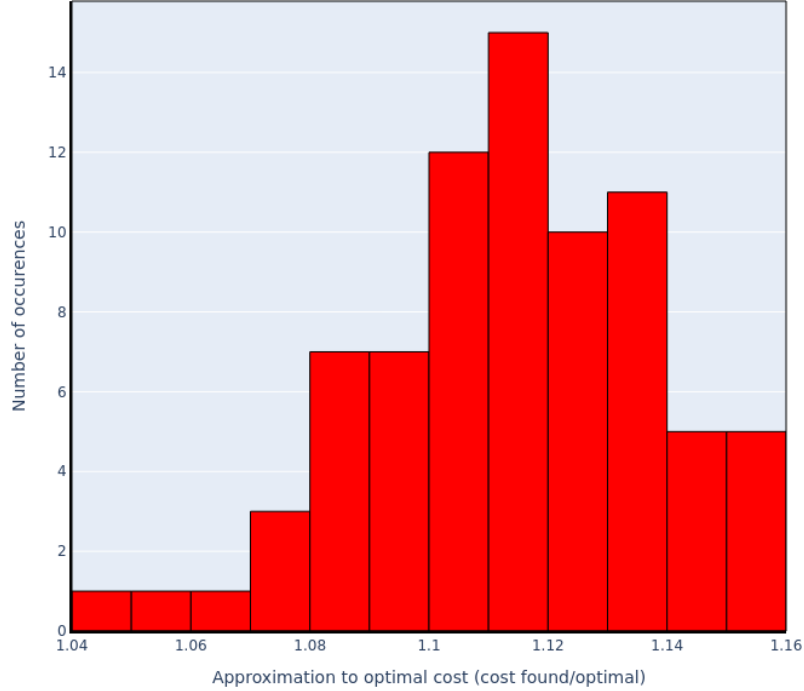


Figure 3. Histogram of approximation rates

As we can see, the Christofides algorithm gives a better answer: usually it goes from 1.08 to 1.16 time the optimal. In fact, it has an maximum distance of the optimal of 1.5 times it. It looks pretty good, but there is some disadvantages, that will be spoken above.

The first one is the time complexity: as in TATT, it makes an MST, which takes $O(n^2)$. However, the next step is more complex than this. Using the same notations than the subsection 2.3, a minimum perfect matching in G is calculated using Blossom algorithm, and this has a complexity of $O(|V(G)|^3)$. So, this depends of how many vertices with an odd degree exists in T , but in general this is the longest step of the Christofides algorithm. The next step, find an Eulerian cycle using Hierholzer's algorithm, can be done in $O(E(H) \times \log_2(E(H)))$. We know that $E(H)$ is at most $n - 1 + \frac{n}{2}$ ($n - 1$ edges from the MST and $\frac{n}{2}$ edges from M), so we can change this complexiy to $O(n \times \log_2 n)$. To finish, the transformation of this cycle in an hamiltonian cycle is done in $O(E(H)) = O(n)$, then the final complexity is $O(n^2 + |V(G)|^3)$.

As we can see in the graphic, the complexity of the algorithm tends to be dominated by the Blossom algorithm (as the factor $O(n^2)$ is also present in TATT algorithm, and it is quite faster than Christofides).

Number of points x execution time using the Christofides algorithm




Figure 4. Relation between the number of points and execution time of Christofides algorithm

Another disadvantage of Christofides algorithm is the space complexity. Here, it is necessary to really create the complete induced subgraph of the vertices with an odd degree in the MST. Let m be the number of these vertices. That takes $O(m^2)$ of space complexity. The Blossom algorithm also takes $O(m^2)$. The last part of the algorithm is cheaper: the creation of the graph H requires $O(E(H))$ (in other words, $O(n)$), and the Hierholzer's algorithm needs the same space. So, the total space complexity of the algorithm is $O(n + m^2)$, and the second term, m^2 , is usually bigger than the term n .

4. Conclusion

As you could see, with a Branch and Bound algorithm, it takes so much time to find the optimal answer, even in small instances. So this algorithm has very specific applications in real life, given that we can not let a computer calculating the answer for days, or even more than this.

However, Christofides and the TATT algorithms have practical applications. Christofides algorithm gives a better answer, so maybe you are thinking that it is always better to use it instead the TATT algorithm. But the Christofides algorithm is way more expansive than TATT, in space and time.

So, with a large set of points/vertices, maybe it's better use the TATT algorithm in order to get a faster answer (unless you have time to wait, and the memory it takes). But, to small or medium set of points, maybe it is worth it to use the Christofides algorithm,

given that the answer is closer to the optimal.

References

- [Dezső et al. 2011] Dezső, B., Jüttner, A., and Kovács, P. (2011). Lemon – an open source c++ graph template library. *Electronic Notes in Theoretical Computer Science*, 264(5):23–45.
- [Deřneko et al. 1997] Deřneko, V. G., Van der Veen, J. A., Rudolf, R., and Woeginger, G. J. (1997). Three easy special cases of the euclidean travelling salesman problem. *RAIRO - Operations Research*, 31(4):343–362.
- [Reinelt 1991] Reinelt, G. (1991). Tspplib—a traveling salesman problem library. *ORSA Journal on Computing*, 3(4):376–384.